

```
In [11]: ▶ import pandas as pd

taxi_owner = pd.read_pickle('taxi_owners.p')
taxi_owner.head()
```

Out[11]:

	rid	vid	owner	address	zip
0	T6285	6285	AGEAN TAXI LLC	4536 N. ELSTON AVE.	60630
1	T4862	4862	MANGIB CORP.	5717 N. WASHTENAW AVE.	60659
2	T1495	1495	FUNRIDE, INC.	3351 W. ADDISON ST.	60618
3	T4231	4231	ALQUSH CORP.	6611 N. CAMPBELL AVE.	60645
4	T5971	5971	EUNIFFORD INC.	3351 W. ADDISON ST.	60618

```
In [7]: ▶ taxi_owner.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3519 entries, 0 to 3518
Data columns (total 5 columns):
#   Column      Non-Null Count  Dtype
---  -
0   rid         3519 non-null   object
1   vid         3519 non-null   object
2   owner       3519 non-null   object
3   address     3519 non-null   object
4   zip         3519 non-null   object
dtypes: object(5)
memory usage: 137.6+ KB
```

```
In [9]: ▶ taxi_owner.describe()
```

Out[9]:

	rid	vid	owner	address	zip
count	3519	3519	3519	3519	3519
unique	3519	3519	2375	317	44
top	T6285	6285	CHICAGO SEVEN INC	3351 W. ADDISON ST.	60618
freq	1	1	21	639	798

```
In [14]: ▶ taxi_owner.shape
```

Out[14]: (3519, 5)

In [16]: `taxi_owner.values`

```
Out[16]: array([[ 'T6285', '6285', 'AGEAN TAXI LLC', '4536 N. ELSTON AVE.',
                '60630'],
               [ 'T4862', '4862', 'MANGIB CORP.', '5717 N. WASHTENAW AVE.',
                '60659'],
               [ 'T1495', '1495', 'FUNRIDE, INC.', '3351 W. ADDISON ST.', '60618'
               ],
               ...,
               [ 'T3465', '3465', 'AMIR EXPRESS INC', '3351 W. ADDISON ST.',
                '60618'],
               [ 'T1962', '1962', 'KARY CAB COMPANY', '4707 N. KENTON AVE.',
                '60630'],
               [ 'T1031', '1031', 'NECT 42 LLC', '6500 N. WESTERN AVE.', '60645'
               ]],
          dtype=object)
```

In [17]: `taxi_owner.columns`

```
Out[17]: Index(['rid', 'vid', 'owner', 'address', 'zip'], dtype='object')
```

In [18]: `taxi_owner.index`

```
Out[18]: RangeIndex(start=0, stop=3519, step=1)
```

In [55]: `homelessnessdf = pd.read_csv('homelessness.csv')`
`homelessnessdf.head()`

```
Out[55]:
```

	Unnamed: 0	region	state	individuals	family_members	state_pop
0	0	East South Central	Alabama	2570.0	864.0	4887681
1	1	Pacific	Alaska	1434.0	582.0	735139
2	2	Mountain	Arizona	7259.0	2606.0	7158024
3	3	West South Central	Arkansas	2280.0	432.0	3009733
4	4	Pacific	California	109008.0	20964.0	39461588

In [29]: `print(homelessnessdf.columns)`

```
Index(['Unnamed: 0', 'region', 'state', 'individuals', 'family_members',
      'state_pop'],
      dtype='object')
```

In []:

In []: `#Ex1`

```
In [33]: homelessness_ind = homelessnessdf.sort_values("individuals", ascending = True)
homelessness_ind.head()
```

Out[33]:

	Unnamed: 0	region	state	individuals	family_members	state_pop
50	50	Mountain	Wyoming	434.0	205.0	577601
34	34	West North Central	North Dakota	467.0	75.0	758080
7	7	South Atlantic	Delaware	708.0	374.0	965479
39	39	New England	Rhode Island	747.0	354.0	1058287
45	45	New England	Vermont	780.0	511.0	624358

```
In [ ]: #Ex2
```

```
In [34]: homelessness_fam = homelessnessdf.sort_values("family_members", ascending = True)
homelessness_fam.head()
```

Out[34]:

	Unnamed: 0	region	state	individuals	family_members	state_pop
32	32	Mid-Atlantic	New York	39827.0	52070.0	19530351
4	4	Pacific	California	109008.0	20964.0	39461588
21	21	New England	Massachusetts	6811.0	13257.0	6882635
9	9	South Atlantic	Florida	21443.0	9587.0	21244317
43	43	West South Central	Texas	19199.0	6111.0	28628666

```
In [ ]: #Ex3
```

```
In [35]: homelessness_reg_fam = homelessnessdf.sort_values(["region", "family_members"])
homelessness_reg_fam.head()
```

Out[35]:

	Unnamed: 0	region	state	individuals	family_members	state_pop
13	13	East North Central	Illinois	6752.0	3891.0	12723071
35	35	East North Central	Ohio	6929.0	3320.0	11676341
22	22	East North Central	Michigan	5209.0	3142.0	9984072
49	49	East North Central	Wisconsin	2740.0	2167.0	5807406
14	14	East North Central	Indiana	3776.0	1482.0	6695497

```
In [ ]: #Ex4
```

```
In [40]: ▶ state_fam = homelessnessdf[["state", "family_members"]]
state_fam.head()
```

Out[40]:

	state	family_members
0	Alabama	864.0
1	Alaska	582.0
2	Arizona	2606.0
3	Arkansas	432.0
4	California	20964.0

```
In [ ]: ▶ #Ex5
```

```
In [49]: ▶ ind_gt_10k = homelessnessdf[homelessnessdf["individuals"] > 10000]
ind_gt_10k.head()
```

Out[49]:

	Unnamed: 0		region	state	individuals	family_members	state_pop
4	4		Pacific	California	109008.0	20964.0	39461588
9	9		South Atlantic	Florida	21443.0	9587.0	21244317
32	32		Mid-Atlantic	New York	39827.0	52070.0	19530351
37	37		Pacific	Oregon	11139.0	3337.0	4181886
43	43		West South Central	Texas	19199.0	6111.0	28628666

```
In [ ]: ▶ #Ex6
```

```
In [48]: ▶ mountain_reg = homelessnessdf[homelessnessdf["region"] == "Mountain"]
mountain_reg.head()
```

Out[48]:

	Unnamed: 0		region	state	individuals	family_members	state_pop
2	2	Mountain	Arizona		7259.0	2606.0	7158024
5	5	Mountain	Colorado		7607.0	3250.0	5691287
12	12	Mountain	Idaho		1297.0	715.0	1750536
26	26	Mountain	Montana		983.0	422.0	1060665
28	28	Mountain	Nevada		7058.0	486.0	3027341

```
In [ ]: ▶ #Ex7
```

```
In [54]: fam_lt_1k_pac = homelessnessdf[(homelessnessdf["family_members"] < 1000) &
fam_lt_1k_pac.head()
```

Out[54]:

	Unnamed: 0	region	state	individuals	family_members	state_pop
1	1	Pacific	Alaska	1434.0	582.0	735139

```
In [56]: #Ex8
```

```
In [58]: south_mid_atlantic = homelessnessdf[(homelessnessdf["region"].isin(["South
south_mid_atlantic
```

Out[58]:

	Unnamed: 0	region	state	individuals	family_members	state_pop
7	7	South Atlantic	Delaware	708.0	374.0	965479
8	8	South Atlantic	District of Columbia	3770.0	3134.0	701547
9	9	South Atlantic	Florida	21443.0	9587.0	21244317
10	10	South Atlantic	Georgia	6943.0	2556.0	10511131
20	20	South Atlantic	Maryland	4914.0	2230.0	6035802
30	30	Mid-Atlantic	New Jersey	6048.0	3350.0	8886025
32	32	Mid-Atlantic	New York	39827.0	52070.0	19530351
33	33	South Atlantic	North Carolina	6451.0	2817.0	10381615
38	38	Mid-Atlantic	Pennsylvania	8163.0	5349.0	12800922
40	40	South Atlantic	South Carolina	3082.0	851.0	5084156
46	46	South Atlantic	Virginia	3928.0	2047.0	8501286
48	48	South Atlantic	West Virginia	1021.0	222.0	1804291

```
In [59]: #Ex9
```

```
In [64]: canu = ["California", "Arizona", "Nevada", "Utah"]

mojave_homelessness = homelessnessdf[homelessnessdf["state"].isin(canu)]
mojave_homelessness.head()
```

Out[64]:

	Unnamed: 0	region	state	individuals	family_members	state_pop
2	2	Mountain	Arizona	7259.0	2606.0	7158024
4	4	Pacific	California	109008.0	20964.0	39461588
28	28	Mountain	Nevada	7058.0	486.0	3027341
44	44	Mountain	Utah	1904.0	972.0	3153550

```
In [65]: #Ex10
```

```
In [66]: ► homelessnessdf["total"] = homelessnessdf["individuals"] + homelessnessdf["state_pop"]
homelessnessdf.head(50)
```


Out[66]:

	Unnamed: 0	region	state	individuals	family_members	state_pop	total
0	0	East South Central	Alabama	2570.0	864.0	4887681	3434.0
1	1	Pacific	Alaska	1434.0	582.0	735139	2016.0
2	2	Mountain	Arizona	7259.0	2606.0	7158024	9865.0
3	3	West South Central	Arkansas	2280.0	432.0	3009733	2712.0
4	4	Pacific	California	109008.0	20964.0	39461588	129972.0
5	5	Mountain	Colorado	7607.0	3250.0	5691287	10857.0
6	6	New England	Connecticut	2280.0	1696.0	3571520	3976.0
7	7	South Atlantic	Delaware	708.0	374.0	965479	1082.0
8	8	South Atlantic	District of Columbia	3770.0	3134.0	701547	6904.0
9	9	South Atlantic	Florida	21443.0	9587.0	21244317	31030.0
10	10	South Atlantic	Georgia	6943.0	2556.0	10511131	9499.0
11	11	Pacific	Hawaii	4131.0	2399.0	1420593	6530.0
12	12	Mountain	Idaho	1297.0	715.0	1750536	2012.0
13	13	East North Central	Illinois	6752.0	3891.0	12723071	10643.0
14	14	East North Central	Indiana	3776.0	1482.0	6695497	5258.0
15	15	West North Central	Iowa	1711.0	1038.0	3148618	2749.0
16	16	West North Central	Kansas	1443.0	773.0	2911359	2216.0
17	17	East South Central	Kentucky	2735.0	953.0	4461153	3688.0
18	18	West South Central	Louisiana	2540.0	519.0	4659690	3059.0
19	19	New England	Maine	1450.0	1066.0	1339057	2516.0

	Unnamed: 0	region	state	individuals	family_members	state_pop	total
20	20	South Atlantic	Maryland	4914.0	2230.0	6035802	7144.0
21	21	New England	Massachusetts	6811.0	13257.0	6882635	20068.0
22	22	East North Central	Michigan	5209.0	3142.0	9984072	8351.0
23	23	West North Central	Minnesota	3993.0	3250.0	5606249	7243.0
24	24	East South Central	Mississippi	1024.0	328.0	2981020	1352.0
25	25	West North Central	Missouri	3776.0	2107.0	6121623	5883.0
26	26	Mountain	Montana	983.0	422.0	1060665	1405.0
27	27	West North Central	Nebraska	1745.0	676.0	1925614	2421.0
28	28	Mountain	Nevada	7058.0	486.0	3027341	7544.0
29	29	New England	New Hampshire	835.0	615.0	1353465	1450.0
30	30	Mid-Atlantic	New Jersey	6048.0	3350.0	8886025	9398.0
31	31	Mountain	New Mexico	1949.0	602.0	2092741	2551.0
32	32	Mid-Atlantic	New York	39827.0	52070.0	19530351	91897.0
33	33	South Atlantic	North Carolina	6451.0	2817.0	10381615	9268.0
34	34	West North Central	North Dakota	467.0	75.0	758080	542.0
35	35	East North Central	Ohio	6929.0	3320.0	11676341	10249.0
36	36	West South Central	Oklahoma	2823.0	1048.0	3940235	3871.0
37	37	Pacific	Oregon	11139.0	3337.0	4181886	14476.0
38	38	Mid-Atlantic	Pennsylvania	8163.0	5349.0	12800922	13512.0
39	39	New England	Rhode Island	747.0	354.0	1058287	1101.0
40	40	South Atlantic	South Carolina	3082.0	851.0	5084156	3933.0

	Unnamed: 0	region	state	individuals	family_members	state_pop	total
41	41	West North Central	South Dakota	836.0	323.0	878698	1159.0
42	42	East South Central	Tennessee	6139.0	1744.0	6771631	7883.0
43	43	West South Central	Texas	19199.0	6111.0	28628666	25310.0
44	44	Mountain	Utah	1904.0	972.0	3153550	2876.0
45	45	New England	Vermont	780.0	511.0	624358	1291.0
46	46	South Atlantic	Virginia	3928.0	2047.0	8501286	5975.0

In [67]:  #Ex11

In [68]:  `homelessnessdf["p_individuals"] = homelessnessdf["individuals"] / homelessnessdf.head()`

Out[68]:

	Unnamed: 0	region	state	individuals	family_members	state_pop	total	p_indiv
0	0	East South Central	Alabama	2570.0	864.0	4887681	3434.0	0.7
1	1	Pacific	Alaska	1434.0	582.0	735139	2016.0	0.7
2	2	Mountain	Arizona	7259.0	2606.0	7158024	9865.0	0.7
3	3	West South Central	Arkansas	2280.0	432.0	3009733	2712.0	0.8
4	4	Pacific	California	109008.0	20964.0	39461588	129972.0	0.8

In []:  #Ex12


```
In [82]: ▶ homelessnessdf["indiv_per_10k"] = 10000 * homelessnessdf['individuals'] / 1000000
high_homelessness = homelessnessdf[homelessnessdf["indiv_per_10k"] > 20]
high_homelessness_srt = high_homelessness.sort_values("indiv_per_10k", ascending=False)
high_homelessness_srt[["state", "indiv_per_10k"]]
```

Out[82]:

	state	indiv_per_10k
8	District of Columbia	53.738381
11	Hawaii	29.079406
4	California	27.623825
37	Oregon	26.636307
28	Nevada	23.314189
47	Washington	21.829195
32	New York	20.392363

In []: ▶

```
In [85]: ▶ def pct40(column):
return column.quantile(0.5)

homelessnessdf["family_members"].agg(pct40)
```

Out[85]: 1482.0

In [89]: ▶

```
sales = pd.read_csv('sales_subset.csv')
sales.head()
```

Out[89]:

	Unnamed: 0	store	type	department	date	weekly_sales	is_holiday	temperature_c
0	0	1	A	1	2010-02-05	24924.50	False	5.727778
1	1	1	A	1	2010-03-05	21827.90	False	8.055556
2	2	1	A	1	2010-04-02	57258.43	False	16.816667
3	3	1	A	1	2010-05-07	17413.94	False	22.527778
4	4	1	A	1	2010-06-04	17558.09	False	27.050000

```
In [92]: # Print the info about the sales DataFrame
print(sales.info())
print()
# Print the mean of weekly_sales
print(sales['weekly_sales'].mean())

# Print the median of weekly_sales
print(sales['weekly_sales'].median())

# Print the maximum of the date column
print(sales['date'].max())

# Print the minimum of the date column
print(sales['date'].min())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10774 entries, 0 to 10773
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Unnamed: 0            10774 non-null  int64
1   store                 10774 non-null  int64
2   type                  10774 non-null  object
3   department            10774 non-null  int64
4   date                  10774 non-null  object
5   weekly_sales          10774 non-null  float64
6   is_holiday            10774 non-null  bool
7   temperature_c         10774 non-null  float64
8   fuel_price_usd_per_l  10774 non-null  float64
9   unemployment          10774 non-null  float64
dtypes: bool(1), float64(4), int64(3), object(2)
memory usage: 768.2+ KB
None

23843.950148505668
12049.064999999999
2012-10-26
2010-02-05
```

```
In [94]: sales_1_1 = sales[(sales['department'] == 1) & (sales['store'] == 1)]

# Sort sales_1_1 by date
sales_1_1 = sales_1_1.sort_values('date', ascending = True)

# Get the cumulative sum of weekly_sales, add as cum_weekly_sales col
sales_1_1['cum_weekly_sales'] = sales_1_1['weekly_sales'].cumsum()

# Get the cumulative max of weekly_sales, add as cum_max_sales col
sales_1_1['cum_max_sales'] = sales_1_1['weekly_sales'].cummax()

# See the columns you calculated
print(sales_1_1[["date", "weekly_sales", "cum_weekly_sales", "cum_max_sales"]])
```

	date	weekly_sales	cum_weekly_sales	cum_max_sales
0	2010-02-05	24924.50	24924.50	24924.50
1	2010-03-05	21827.90	46752.40	24924.50
2	2010-04-02	57258.43	104010.83	57258.43
3	2010-05-07	17413.94	121424.77	57258.43
4	2010-06-04	17558.09	138982.86	57258.43
5	2010-07-02	16333.14	155316.00	57258.43
6	2010-08-06	17508.41	172824.41	57258.43
7	2010-09-03	16241.78	189066.19	57258.43
8	2010-10-01	20094.19	209160.38	57258.43
9	2010-11-05	34238.88	243399.26	57258.43
10	2010-12-03	22517.56	265916.82	57258.43
11	2011-01-07	15984.24	281901.06	57258.43

```
In [95]: #13
```

```
In [96]: store_types = sales.drop_duplicates(subset = ["store", "type"])
store_types.head()
```

Out[96]:

	Unnamed: 0	store	type	department	date	weekly_sales	is_holiday	temperature
0	0	1	A	1	2010-02-05	24924.50	False	5.7277
901	901	2	A	1	2010-02-05	35034.06	False	4.5500
1798	1798	4	A	1	2010-02-05	38724.42	False	6.5333
2699	2699	6	A	1	2010-02-05	25619.00	False	4.6833
3593	3593	10	B	1	2010-02-05	40212.84	False	12.4111

```
In [98]: store_depts = sales.drop_duplicates(subset = ["store", "department"])
store_depts.head()
```

Out[98]:

	Unnamed: 0	store	type	department	date	weekly_sales	is_holiday	temperature_c
0	0	1	A	1	2010-02-05	24924.50	False	5.727778
12	12	1	A	2	2010-02-05	50605.27	False	5.727778
24	24	1	A	3	2010-02-05	13740.12	False	5.727778
36	36	1	A	4	2010-02-05	39954.04	False	5.727778
48	48	1	A	5	2010-02-05	32229.38	False	5.727778

```
In [118]: holiday_dates = sales[sales["is_holiday"] == True].drop_duplicates(subset =
holiday_dates['date']
```

Out[118]:

498	2010-09-10
691	2011-11-25
2315	2010-02-12
6735	2012-09-07
6810	2010-12-31
6815	2012-02-10
6820	2011-09-09

Name: date, dtype: object

```
In [109]: #Ex14
```

```
In [157]: store_types = sales.drop_duplicates(subset = ['store', 'type'])
store_counts = store_types['type'].value_counts()
print(store_counts)

store_props = store_types['type'].value_counts(normalize= True)
print(store_props)

store_depts = sales.drop_duplicates(subset = ['store', 'department'])
dept_counts_sorted = store_depts [ 'department'].value_counts(sort=True)
print(dept_counts_sorted)

dept_props_sorted = store_depts['department'].value_counts(sort = True, no
print(dept_props_sorted)
```

```
A    11
B     1
Name: type, dtype: int64
A    0.916667
B    0.083333
Name: type, dtype: float64
1     12
55    12
72    12
71    12
67    12
..
37    10
48     8
50     6
39     4
43     2
Name: department, Length: 80, dtype: int64
1     0.012917
55    0.012917
72    0.012917
71    0.012917
67    0.012917
...
37    0.010764
48    0.008611
50    0.006459
39    0.004306
43    0.002153
Name: department, Length: 80, dtype: float64
```

In []: ▶

In []: ▶

In []: ▶

In []: ▶ *#Practise*

```
In [145]: ▶ # Calc total weekly sales
sales_all = sales["weekly_sales"].sum()

# Subset for type A stores, calc total weekly sales
sales_A = sales[sales["type"] == "A"]["weekly_sales"].sum()

# Subset for type B stores, calc total weekly sales
sales_B = sales[sales["type"] == "B"]["weekly_sales"].sum()

# Subset for type C stores, calc total weekly sales
sales_C = sales[sales["type"] == "C"]["weekly_sales"].sum()

# Get proportion for each type
sales_propn_by_type = [sales_A, sales_B, sales_C] / sales_all
print(sales_propn_by_type)
```

```
[0.9097747 0.0902253 0.      ]
```

```
In [ ]: ▶ #15
The code first calculates the sum of all weekly sales using the sum() func
Then, it subsets the data for each store type (A, B, and C) and calculates
Finally, it computes the proportion of sales for each store type by dividin
The results are stored in the variable sales_propn_by_type as a list.
```

```
In [ ]: ▶ #Practise
```

```
In [139]: # Import numpy with the alias np
import numpy as np

# For each store type, aggregate weekly_sales: get min, max, mean, and median
sales_stats = sales.groupby('type')['weekly_sales'].agg([min, max, np.mean, np.median])

# Print sales_stats
print(sales_stats)

# For each store type, aggregate unemployment and fuel_price_usd_per_l: get min, max, mean, and median
unemp_fuel_stats = sales.groupby('type')[['unemployment', 'fuel_price_usd_per_l']].agg([min, max, np.mean, np.median])

# Print unemp_fuel_stats
print(unemp_fuel_stats)
```

	min	max	mean	median
type				
A	-1098.0	293966.05	23674.667242	11943.92
B	-798.0	232558.51	25696.678370	13336.08

	unemployment	fuel_price_usd_per_l				
\						
	min	max	mean	median	min	max
type						
A	3.879	8.992	7.972611	8.067	0.664129	1.107410
B	7.170	9.765	9.279323	9.199	0.760023	1.107674

	mean	median
type		
A	0.744619	0.735455
B	0.805858	0.803348

```
In [141]: # Print mean weekly_sales by department and type; fill missing values with 0
print(sales.pivot_table(values = 'weekly_sales', index = 'department', columns = 'type', fill_value = 0))
```

type	A	B
department		
1	30961.725379	44050.626667
2	67600.158788	112958.526667
3	17160.002955	30580.655000
4	44285.399091	51219.654167
5	34821.011364	63236.875000
...
95	123933.787121	77082.102500
96	21367.042857	9528.538333
97	28471.266970	5828.873333
98	12875.423182	217.428333
99	379.123659	0.000000

[80 rows x 2 columns]

```

In [143]: temperatures = pd.read_csv('temperatures.csv')

# Look at temperatures
print(temperatures)
# Set the index of temperatures to city
temperatures_ind = temperatures.set_index('city')

# Look at temperatures_ind
print(temperatures_ind)

# Reset the temperatures_ind index, keeping its contents
print(temperatures_ind.reset_index())

# Reset the temperatures_ind index, dropping its contents
print(temperatures_ind.reset_index(drop = True))
# Make a list of cities to subset on
cities = ["Moscow", "Saint Petersburg"]

# Subset temperatures using square brackets
print(temperatures[temperatures['city'].isin(cities)])

# Subset temperatures_ind using .loc[]
print(temperatures_ind.loc[cities])

```

	Unnamed: 0	date	city	country	avg_temp_c
0	0	2000-01-01	Abidjan	Côte D'Ivoire	27.293
1	1	2000-02-01	Abidjan	Côte D'Ivoire	27.685
2	2	2000-03-01	Abidjan	Côte D'Ivoire	29.061
3	3	2000-04-01	Abidjan	Côte D'Ivoire	28.162
4	4	2000-05-01	Abidjan	Côte D'Ivoire	27.547
...
16495	16495	2013-05-01	Xian	China	18.979
16496	16496	2013-06-01	Xian	China	23.522
16497	16497	2013-07-01	Xian	China	25.251
16498	16498	2013-08-01	Xian	China	24.528
16499	16499	2013-09-01	Xian	China	NaN

[16500 rows x 5 columns]

	Unnamed: 0	date	country	avg_temp_c
city				
Abidjan	0	2000-01-01	Côte D'Ivoire	27.293
Abidjan	1	2000-02-01	Côte D'Ivoire	27.685
Abidjan	2	2000-03-01	Côte D'Ivoire	29.061
Abidjan	3	2000-04-01	Côte D'Ivoire	28.162
Abidjan	4	2000-05-01	Côte D'Ivoire	27.547
...
Xian	16495	2013-05-01	China	18.979
Xian	16496	2013-06-01	China	23.522
Xian	16497	2013-07-01	China	25.251
Xian	16498	2013-08-01	China	24.528
Xian	16499	2013-09-01	China	NaN

[16500 rows x 4 columns]

	city	Unnamed: 0	date	country	avg_temp_c
0	Abidjan	0	2000-01-01	Côte D'Ivoire	27.293

1	Abidjan	1	2000-02-01	Côte D'Ivoire	27.685
2	Abidjan	2	2000-03-01	Côte D'Ivoire	29.061
3	Abidjan	3	2000-04-01	Côte D'Ivoire	28.162
4	Abidjan	4	2000-05-01	Côte D'Ivoire	27.547
...
16495	Xian	16495	2013-05-01	China	18.979
16496	Xian	16496	2013-06-01	China	23.522
16497	Xian	16497	2013-07-01	China	25.251
16498	Xian	16498	2013-08-01	China	24.528
16499	Xian	16499	2013-09-01	China	NaN

[16500 rows x 5 columns]

	Unnamed: 0	date	country	avg_temp_c
0	0	2000-01-01	Côte D'Ivoire	27.293
1	1	2000-02-01	Côte D'Ivoire	27.685
2	2	2000-03-01	Côte D'Ivoire	29.061
3	3	2000-04-01	Côte D'Ivoire	28.162
4	4	2000-05-01	Côte D'Ivoire	27.547
...
16495	16495	2013-05-01	China	18.979
16496	16496	2013-06-01	China	23.522
16497	16497	2013-07-01	China	25.251
16498	16498	2013-08-01	China	24.528
16499	16499	2013-09-01	China	NaN

[16500 rows x 4 columns]


	Unnamed: 0	date	city	country	avg_temp_c
10725	10725	2000-01-01	Moscow	Russia	-7.313
10726	10726	2000-02-01	Moscow	Russia	-3.551
10727	10727	2000-03-01	Moscow	Russia	-1.661
10728	10728	2000-04-01	Moscow	Russia	10.096
10729	10729	2000-05-01	Moscow	Russia	10.357
...
13360	13360	2013-05-01	Saint Petersburg	Russia	12.355
13361	13361	2013-06-01	Saint Petersburg	Russia	17.185
13362	13362	2013-07-01	Saint Petersburg	Russia	17.234
13363	13363	2013-08-01	Saint Petersburg	Russia	17.153
13364	13364	2013-09-01	Saint Petersburg	Russia	NaN

[330 rows x 5 columns]

	Unnamed: 0	date	country	avg_temp_c
city				
Moscow	10725	2000-01-01	Russia	-7.313
Moscow	10726	2000-02-01	Russia	-3.551
Moscow	10727	2000-03-01	Russia	-1.661
Moscow	10728	2000-04-01	Russia	10.096
Moscow	10729	2000-05-01	Russia	10.357
...
Saint Petersburg	13360	2013-05-01	Russia	12.355
Saint Petersburg	13361	2013-06-01	Russia	17.185
Saint Petersburg	13362	2013-07-01	Russia	17.234
Saint Petersburg	13363	2013-08-01	Russia	17.153
Saint Petersburg	13364	2013-09-01	Russia	NaN

[330 rows x 4 columns]

In [144]:  #16

```
In [150]:  temperatures_ind = temperatures.set_index(["country", "city"])
rows_to_keep = [("Brazil", "Rio De Janeiro") , ("Pakistan" , "Lahore")]
subset = temperatures_ind.loc[rows_to_keep]

subset
```

Out[150]:

		Unnamed: 0	date	avg_temp_c
country	city			
Brazil	Rio De Janeiro	12540	2000-01-01	25.974
	Rio De Janeiro	12541	2000-02-01	26.699
	Rio De Janeiro	12542	2000-03-01	26.270
	Rio De Janeiro	12543	2000-04-01	25.750
	Rio De Janeiro	12544	2000-05-01	24.356
...
Pakistan	Lahore	8575	2013-05-01	33.457
	Lahore	8576	2013-06-01	34.456
	Lahore	8577	2013-07-01	33.279
	Lahore	8578	2013-08-01	31.511
	Lahore	8579	2013-09-01	NaN

330 rows × 3 columns

In []: 