

```
In [40]: import pandas as pd
cc_apps = pd.read_csv("cc_approvals.data", header = None)
cc_apps.head()
```

Out[40]:

	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
0	b	30.83	0.000	u	g	w	v	1.25	t	t	1	f	g	00202	0	+
1	a	58.67	4.460	u	g	q	h	3.04	t	t	6	f	g	00043	560	+
2	a	24.50	0.500	u	g	q	h	1.50	t	f	0	f	g	00280	824	+
3	b	27.83	1.540	u	g	w	v	3.75	t	t	5	t	g	00100	3	+
4	b	20.17	5.625	u	g	w	v	1.71	t	f	0	f	s	00120	0	+

```
In [41]: print(cc_apps.describe())
print('\n')

print(cc_apps.info())
print('\n')

cc_apps.tail(17) # or cc_apps.sample()
```

	2	7	10	14
count	690.000000	690.000000	690.000000	690.000000
mean	4.758725	2.223406	2.400000	1017.385507
std	4.978163	3.346513	4.86294	5210.102598
min	0.000000	0.000000	0.000000	0.000000
25%	1.000000	0.165000	0.000000	0.000000
50%	2.750000	1.000000	0.000000	5.000000
75%	7.207500	2.625000	3.000000	395.500000
max	28.000000	28.500000	67.000000	100000.000000

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 690 entries, 0 to 689
Data columns (total 16 columns):
#   Column  Non-Null Count  Dtype
---  -
0    0      690 non-null      object
1    1      690 non-null      object
2    2      690 non-null      float64
3    3      690 non-null      object
4    4      690 non-null      object
5    5      690 non-null      object
6    6      690 non-null      object
7    7      690 non-null      float64
8    8      690 non-null      object
9    9      690 non-null      object
10   10     690 non-null      int64
11   11     690 non-null      object
12   12     690 non-null      object
13   13     690 non-null      object
14   14     690 non-null      int64
15   15     690 non-null      object
dtypes: float64(2), int64(2), object(12)
memory usage: 86.4+ KB
None
```

Out[41]:

	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
<b>673</b>	?	29.50	2.000	y	p	e	h	2.000	f	f	0	f	g	00256	17	-
<b>674</b>	a	37.33	2.500	u	g	i	h	0.210	f	f	0	f	g	00260	246	-
<b>675</b>	a	41.58	1.040	u	g	aa	v	0.665	f	f	0	f	g	00240	237	-
<b>676</b>	a	30.58	10.665	u	g	q	h	0.085	f	t	12	t	g	00129	3	-

	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
677	b	19.42	7.250	u	g	m	v	0.040	f	t	1	f	g	00100	1	-
678	a	17.92	10.210	u	g	ff	ff	0.000	f	f	0	f	g	00000	50	-
679	a	20.08	1.250	u	g	c	v	0.000	f	f	0	f	g	00000	0	-
680	b	19.50	0.290	u	g	k	v	0.290	f	f	0	f	g	00280	364	-
681	b	27.83	1.000	y	p	d	h	3.000	f	f	0	f	g	00176	537	-
682	b	17.08	3.290	u	g	i	v	0.335	f	f	0	t	g	00140	2	-
683	b	36.42	0.750	y	p	d	v	0.585	f	f	0	f	g	00240	3	-
684	b	40.58	3.290	u	g	m	v	3.500	f	f	0	t	s	00400	0	-
685	b	21.08	10.085	y	p	e	h	1.250	f	f	0	f	g	00260	0	-
686	a	22.67	0.750	u	g	c	v	2.000	f	t	2	t	g	00200	394	-
687	a	25.25	13.500	y	p	ff	ff	2.000	f	t	1	t	g	00200	1	-
688	b	17.92	0.205	u	a	aa	v	0.040	f	f	0	f	a	00280	750	-

```
In [42]: from sklearn.model_selection import train_test_split

print(cc_apps.corr())

#Drop the features 11 and 13
cc_apps = cc_apps.drop([11, 13], axis=1)

# Split into train and test sets
cc_apps_train, cc_apps_test = train_test_split(cc_apps, test_size=0.33, random
```

	2	7	10	14
2	1.000000	0.298902	0.271207	0.123121
7	0.298902	1.000000	0.322330	0.051345
10	0.271207	0.322330	1.000000	0.063692
14	0.123121	0.051345	0.063692	1.000000

```
In [43]: # Import numpy
import numpy as np

# Replace the '?'s with NaN in the train and test sets
cc_apps_train = cc_apps_train.replace('?', np.NaN)
cc_apps_test = cc_apps_test.replace('?', np.NaN)
```

```
In [44]: # Impute the missing values with mean imputation
cc_apps_train.fillna(cc_apps_train.mean(), inplace=True)
cc_apps_test.fillna(cc_apps_train.mean(), inplace=True)

# Count the number of NaNs in the datasets and print the counts to verify
print(cc_apps_train.isnull().sum())
print(cc_apps_test.isnull().sum())
```

```
0      8
1      5
2      0
3      6
4      6
5      7
6      7
7      0
8      0
9      0
10     0
12     0
14     0
15     0
dtype: int64
0      4
1      7
2      0
3      0
4      0
5      2
6      2
7      0
8      0
9      0
10     0
12     0
14     0
15     0
dtype: int64
```

C:\Users\Lut Lat Aung\AppData\Local\Temp\ipykernel\_29180\3580017964.py:2: FutureWarning: Dropping of nuisance columns in DataFrame reductions (with 'numeric\_only=None') is deprecated; in a future version this will raise TypeError. Select only valid columns before calling the reduction.

```
cc_apps_train.fillna(cc_apps_train.mean(), inplace=True)
```

C:\Users\Lut Lat Aung\AppData\Local\Temp\ipykernel\_29180\3580017964.py:3: FutureWarning: Dropping of nuisance columns in DataFrame reductions (with 'numeric\_only=None') is deprecated; in a future version this will raise TypeError. Select only valid columns before calling the reduction.

```
cc_apps_test.fillna(cc_apps_train.mean(), inplace=True)
```

```
In [45]: for col in cc_apps_train.columns: # Iterate over each column of cc_apps_train

        if cc_apps_train[col].dtypes == 'object': # Check if the column is of object type
            # Impute with the most frequent value
            # The value_counts() function returns a Series that contain counts of values in descending order so that its first element will be the most frequent
            cc_apps_train = cc_apps_train.fillna(cc_apps_train[col].value_counts().index[0])
            cc_apps_test = cc_apps_test.fillna(cc_apps_train[col].value_counts().index[0])

# Count the number of NaNs in the dataset and print the counts to verify
print(cc_apps_train.isnull().sum())
print(cc_apps_test.isnull().sum())
# At this point, there is no missing values.
```

```
0      0
1      0
2      0
3      0
4      0
5      0
6      0
7      0
8      0
9      0
10     0
12     0
14     0
15     0
dtype: int64
0      0
1      0
2      0
3      0
4      0
5      0
6      0
7      0
8      0
9      0
10     0
12     0
14     0
15     0
dtype: int64
```

```
In [46]: # Convert the categorical features in the train and test sets independently
print(cc_apps_train)
cc_apps_train = pd.get_dummies(cc_apps_train)
cc_apps_test = pd.get_dummies(cc_apps_test)
print(cc_apps_train)
# Reindex the columns of the test set aligning with the train set
cc_apps_test = cc_apps_test.reindex(columns=cc_apps_train.columns, fill_value=
```

	0	1	2	3	4	5	6	7	8	9	10	12	14	15
382	a	24.33	2.500	y	p	i	bb	4.500	f	f	0	g	456	-
137	b	33.58	2.750	u	g	m	v	4.250	t	t	6	g	0	+
346	b	32.25	1.500	u	g	c	v	0.250	f	f	0	g	122	-
326	b	30.17	1.085	y	p	c	v	0.040	f	f	0	g	179	-
33	a	36.75	5.125	u	g	e	v	5.000	t	f	0	g	4000	+
..	..	...	...	..	..	..	..	...	..	..	..	..	...	..
71	b	34.83	4.000	u	g	d	bb	12.500	t	f	0	g	0	-
106	b	28.75	1.165	u	g	k	v	0.500	t	f	0	s	0	-
270	b	37.58	0.000	b	b	b	b	0.000	f	f	0	p	0	+
435	b	19.00	0.000	y	p	ff	ff	0.000	f	t	4	g	1	-
102	b	18.67	5.000	u	g	q	v	0.375	t	t	2	g	38	-

[462 rows x 14 columns]

	2	7	10	14	0_a	0_b	1_13.75	1_15.83	1_15.92	1_16.00	\
382	2.500	4.500	0	456	1	0	0	0	0	0	
137	2.750	4.250	6	0	0	1	0	0	0	0	
346	1.500	0.250	0	122	0	1	0	0	0	0	
326	1.085	0.040	0	179	0	1	0	0	0	0	
33	5.125	5.000	0	4000	1	0	0	0	0	0	
..	...	...	..	...	...	...	...	...	...	...	
71	4.000	12.500	0	0	0	1	0	0	0	0	
106	1.165	0.500	0	0	0	1	0	0	0	0	
270	0.000	0.000	0	0	0	1	0	0	0	0	
435	0.000	0.000	4	1	0	1	0	0	0	0	
102	5.000	0.375	2	38	0	1	0	0	0	0	

	...	6_z	8_f	8_t	9_f	9_t	12_g	12_p	12_s	15_+	15_-
382	...	0	1	0	1	0	1	0	0	0	1
137	...	0	0	1	0	1	1	0	0	1	0
346	...	0	1	0	1	0	1	0	0	0	1
326	...	0	1	0	1	0	1	0	0	0	1
33	...	0	0	1	1	0	1	0	0	1	0
..	...	...	...	...	...	...	...	...	...	...	...
71	...	0	0	1	1	0	1	0	0	0	1
106	...	0	0	1	1	0	0	0	1	0	1
270	...	0	1	0	1	0	0	1	0	1	0
435	...	0	1	0	0	1	1	0	0	0	1
102	...	0	0	1	0	1	1	0	0	0	1

[462 rows x 334 columns]

```
In [47]: # Import MinMaxScaler
from sklearn.preprocessing import MinMaxScaler

# Segregate features and labels into separate variables
X_train, y_train = cc_apps_train.iloc[:, :-1].values, cc_apps_train.iloc[:, [-1]]
X_test, y_test = cc_apps_test.iloc[:, :-1].values, cc_apps_test.iloc[:, [-1]].

# Instantiate MinMaxScaler and use it to rescale X_train and X_test
scaler = MinMaxScaler(feature_range=(0, 1))
rescaledX_train = scaler.fit_transform(X_train)
rescaledX_test = scaler.transform(X_test)
```

```
In [48]: # Import LogisticRegression
from sklearn.linear_model import LogisticRegression

# Instantiate a LogisticRegression classifier with default parameter values
logreg = LogisticRegression()

# Fit logreg to the train set
logreg.fit(rescaledX_train, y_train)
```

```
C:\Users\Lut Lat Aung\anaconda3\lib\site-packages\sklearn\utils\validation.p
y:993: DataConversionWarning: A column-vector y was passed when a 1d array wa
s expected. Please change the shape of y to (n_samples, ), for example using
ravel().
y = column_or_1d(y, warn=True)
```

```
Out[48]: LogisticRegression()
```

```
In [49]: # Import confusion_matrix
from sklearn.metrics import confusion_matrix

# Use logreg to predict instances from the test set and store it
y_pred = logreg.predict(rescaledX_test)

# Get the accuracy score of Logreg model and print it
print("Accuracy of logistic regression classifier: ", logreg.score(rescaledX_t

# Print the confusion matrix of the Logreg model
confusion_matrix(y_test, y_pred)
```

```
Accuracy of logistic regression classifier: 1.0
```

```
Out[49]: array([[103,  0],
               [ 0, 125]], dtype=int64)
```

In [50]:

```
Accuracy of the best KNN model: 0.7368421052631579
```

```
C:\Users\Lut Lat Aung\anaconda3\lib\site-packages\sklearn\neighbors\_classification.py:198: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().
    return self._fit(X, y)
```



In [66]: *#Task 1*

```
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsRegressor
from sklearn.metrics import mean_squared_error
from sklearn.model_selection import GridSearchCV
from sklearn.preprocessing import MinMaxScaler, LabelEncoder

# Load data
cc_apps = pd.read_csv("cc_approvals.data", header=None)

# Data Preprocessing
# Handling Missing Values (replace '?' with NaN)
cc_apps.replace('?', np.nan, inplace=True)
cc_apps.fillna(cc_apps.median(), inplace=True) # Filling missing values with

# Convert non-numeric columns to strings
for column in cc_apps.columns:
    if cc_apps[column].dtype == 'object':
        cc_apps[column] = cc_apps[column].astype(str)

# Label Encoding for non-numeric columns
label_encoders = {}
for column in cc_apps.columns:
    if cc_apps[column].dtype == 'object':
        label_encoders[column] = LabelEncoder()
        cc_apps[column] = label_encoders[column].fit_transform(cc_apps[column])

# Drop features 11 and 13
cc_apps = cc_apps.drop([11, 13], axis=1)

# Split data into features and target variable
X = cc_apps.drop(columns=[2]) # Replace 'target_column' with your actual targ
y = cc_apps[2]

# Split into train and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33, rand

# Find the best k using Grid Search for KNeighborsRegressor
param_grid = {'n_neighbors': range(1, 21)} # Trying k values from 1 to 20
knn = KNeighborsRegressor()
grid_search = GridSearchCV(knn, param_grid, cv=5)
grid_search.fit(X_train_scaled, y_train)

# Print the best k value
print("Best K: ", grid_search.best_params_['n_neighbors'])

# Train the final KNN model with the best k value
best_k = grid_search.best_params_['n_neighbors']
knn = KNeighborsRegressor(n_neighbors=best_k)
knn.fit(X_train_scaled, y_train)
```

```
test_accuracy = knn.score(X_test, y_test)

print("Accuracy of the best KNN model: ", test_accuracy)
```

C:\Users\Lut Lat Aung\AppData\Local\Temp\ipykernel\_29180\863855895.py:17: FutureWarning: Dropping of nuisance columns in DataFrame reductions (with 'numeric\_only=None') is deprecated; in a future version this will raise TypeError. Select only valid columns before calling the reduction.

```
cc_apps.fillna(cc_apps.median(), inplace=True) # Filling missing values with median for numerical columns
```

Best K: 18

Accuracy of the best KNN model: -0.43739902524631247

In [ ]:

In [ ]:

```
In [55]: import pandas as pd
cc_apps = pd.read_csv("cc_approvals.data", header = None)
print(cc_apps.head())
print(cc_apps.describe())
print('\n')
print(cc_apps.info())
print('\n')
cc_apps.tail(17)
cc_apps.fillna(cc_apps.mean(), inplace=True)
print(cc_apps.isnull().sum())
import numpy as np
cc_apps = cc_apps.replace('?', np.NaN)
for col in cc_apps.columns:

    if cc_apps[col].dtypes == 'object':

        cc_apps = cc_apps.fillna(cc_apps[col].value_counts().index[0])
print(cc_apps.isnull().sum())
print(cc_apps.describe())
print('\n')
print(cc_apps.info())
print('\n')
cc_apps.tail(17)
cc_apps = pd.get_dummies(cc_apps)
print(cc_apps)

from sklearn.preprocessing import MinMaxScaler

X_train, y_train = cc_apps.iloc[:, :-1].values, cc_apps.iloc[:, [-1]].values
X_test, y_test = cc_apps.iloc[:, :-1].values, cc_apps.iloc[:, [-1]].values
scaler = MinMaxScaler(feature_range=(0, 1))
rescaledX_train = scaler.fit_transform(X_train)
rescaledX_test = scaler.transform(X_test)

from sklearn.linear_model import LogisticRegression

logreg = LogisticRegression()
logreg.fit(rescaledX_train,y_train)

from sklearn.model_selection import train_test_split

print(cc_apps.corr())
cc_apps_train, cc_apps_test = train_test_split(cc_apps, test_size=0.33, random

from sklearn.metrics import confusion_matrix

y_pred = logreg.predict(rescaledX_test)
print("Accuracy of logistic regression classifier: ", logreg.score(rescaledX_t
confusion_matrix(y_test,y_pred)

from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score

for k in range(1, 10):
    knn_classifier = KNeighborsClassifier(n_neighbors=k)
```

```

knn_classifier.fit(X_train, y_train)
y_pred = knn_classifier.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
print(f'k = {k}: Accuracy = {accuracy:.2f}')

```

	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
0	b	30.83	0.000	u	g	w	v	1.25	t	t	1	f	g	00202	0	+
1	a	58.67	4.460	u	g	q	h	3.04	t	t	6	f	g	00043	560	+
2	a	24.50	0.500	u	g	q	h	1.50	t	f	0	f	g	00280	824	+
3	b	27.83	1.540	u	g	w	v	3.75	t	t	5	t	g	00100	3	+
4	b	20.17	5.625	u	g	w	v	1.71	t	f	0	f	s	00120	0	+
			2				7				10				14	
count		690.000000		690.000000				690.000000				690.000000				
mean		4.758725		2.223406				2.400000				1017.385507				
std		4.978163		3.346513				4.86294				5210.102598				
min		0.000000		0.000000				0.000000				0.000000				
25%		1.000000		0.165000				0.000000				0.000000				
50%		2.750000		1.000000				0.000000				5.000000				
75%		7.207500		2.625000				3.000000				395.500000				
max		28.000000		28.500000				67.000000				100000.000000				

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 690 entries, 0 to 689
Data columns (total 16 columns):

```

In [ ]: