

# Lesson 10

# Artificial Neural Networks

Mathematics and Statistics for Data Science

Tapanan Yeophantong

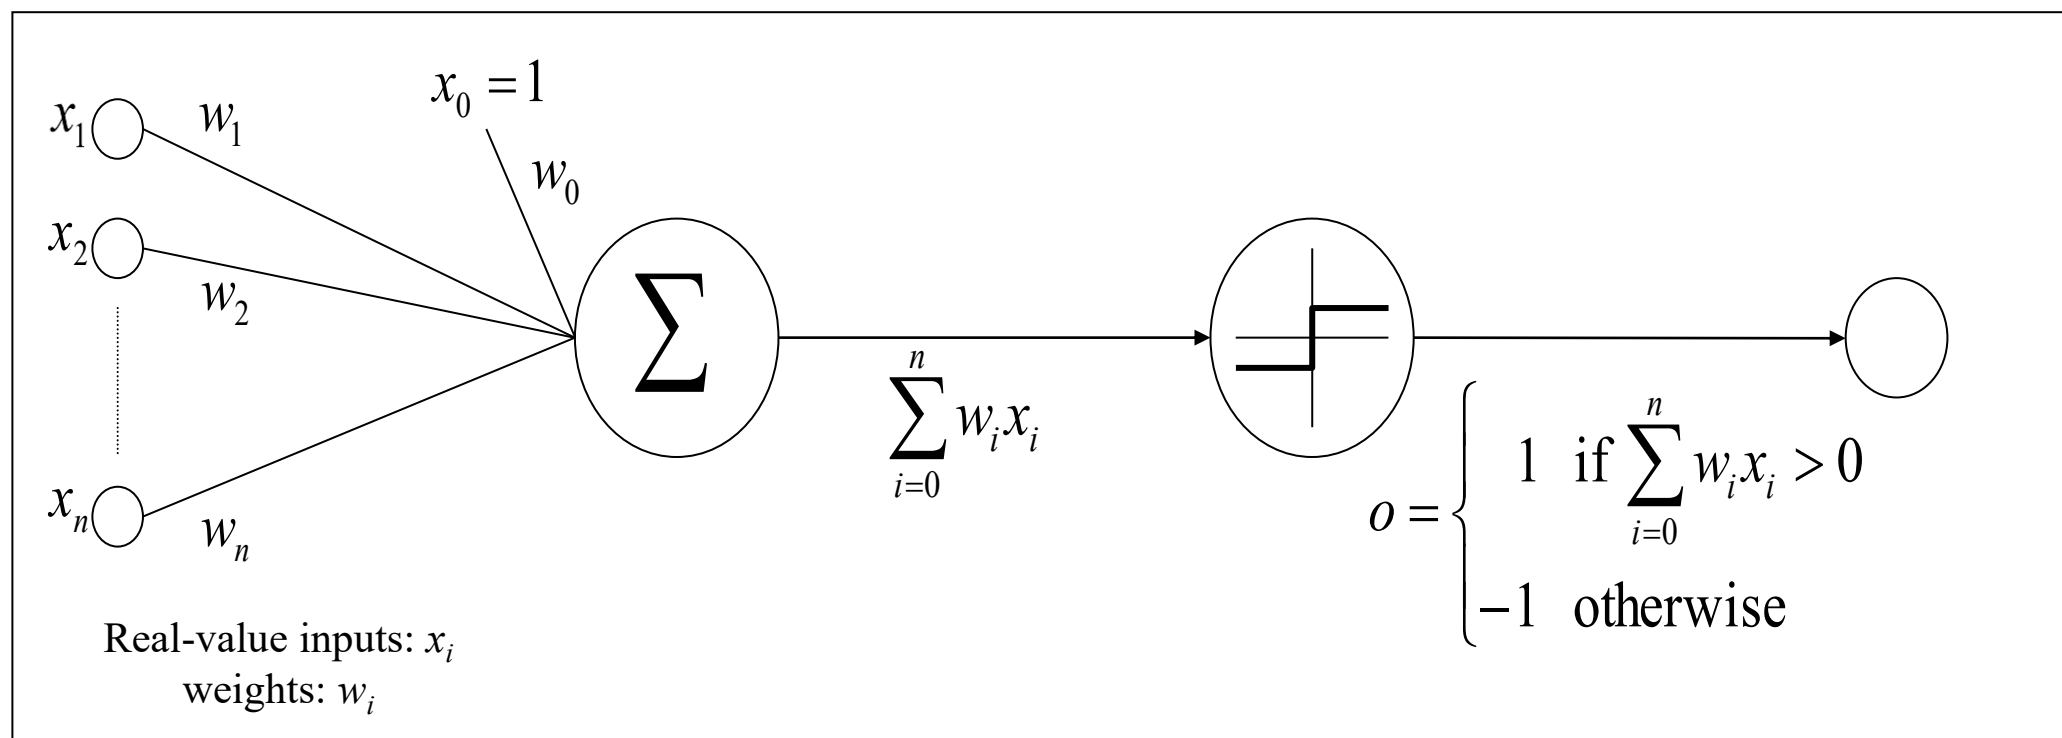
Vincent Mary School of Science and Technology

Assumption University

# Content


- Perceptrons
- Stochastic gradient descent algorithm
- Backpropagation & multi-layered networks
- Deep learning architectures & applications

# Perceptrons



# Representational Power


- Decision surfaces represented by a two-input perceptron



- Perceptron outputs a 1 for instances on one side of hyperplane and outputs a -1 for instances on other side.

# Non-separability Problem

- Some sets of positive and negative examples cannot be separated by a perceptron.



- Those that can be separated by a hyperplane are linearly separable examples.

# Multilayer Perceptrons

- Multilayer perceptron can compute all boolean functions.
- In fact, every Boolean function can be represented by a network of interconnected units.
- Only two-level depth is sufficient.

# Training a Perceptron

- Learning the weights of a single perceptron.
- Produce correct +1 output given each training example
- Two algorithms:
  - Perceptron Rule
  - Delta Rule
- Guaranteed to converge to different hypotheses.

# Perceptron Training Rule

- Begin with random weights.
- Iteratively apply the perceptron to each training example.
- Modify the weights whenever it misclassifies a sample.
- Iterate as many times as needed until all training samples are correctly classified.



# Perceptron Training Rule

- Training involves changing weight  $w_i$  associated with  $x_i$ 
  - Weight is revised by
$$w_i \leftarrow w_i + \Delta w_i$$
  - where  $\Delta w_i = \eta(t - o)x_i$
  - Learning rate  $\eta$  is a positive constant that moderates the degree to which weights are changed at each step. It is set to some small value and sometimes decays as the number of iterations increases.
  - $t$  is the target output.
  - $o$  is the output generated by the perceptron.

# Perceptron Rule Convergence

- Suppose training example is correctly classified
  - $(t - o) = 0$  so no weights updated
- But if perceptron outputs  $-1$  when target output is  $+1$ 
  - weights must be altered to increase value of  $w.x$
  - for all  $i$ , if  $x_i > 0$  (all attributes are positive), increasing  $w_i$  will bring perceptron closer to correctly classifying the example.
  - training rule increases  $w_i$  because  $(t - o)$ ,  $\eta$  and  $x_i$  are all positive.

# Perceptron Convergence

- Perceptron learning procedure is guarantee to converge provided that:
  - The training samples are linearly separable, and
  - A sufficiently small  $\eta$  is used.

# Exercise 1

- **Goal:** Try a perceptron on scikit-learn.
- **Instructions:**
  - Create datasets corresponding to ADD, OR and XOR.
  - Train a perceptron with each dataset.
  - Check your results.

# Content

- Perceptrons
- Stochastic gradient descent algorithm
- Backpropagation & multi-layered networks
- Deep learning architectures & applications

# Delta Rule

- Perceptron rule fails to converge if examples are not linearly separable.
- Delta rule is designed to overcome this difficulty.
- If samples are not linearly separable, Delta rule converges toward a **best-fit approximation** of target concept.
- The key idea is to use **gradient descent** to search hypothesis space of possible weight vectors.


# Overview of the Delta Rule

- Unthresholded perceptron:  $o(x) = w \cdot x$
- Training error of a hypothesis weight vectors:

$$E(\vec{w}) \equiv \frac{1}{2} \sum_{d \in D} (t_d - o_d)^2$$

- where  $D$  is the set of training samples
- $t_d$  is target output for training sample  $d$
- $o_d$  is output for training sample  $d$

# Visualising Hypothesis Space





# Gradient Descent

- Gradient descent determines a weight vector that minimises  $E$ 
  - Starting with an arbitrary initial weight vector.
  - Repeatedly modifying it in small steps.
  - At each step, weight vector is modified in the direction that produces the steepest descent along the error surface.

# Weight Update Rule

- Pick an initial random weight vector.
- Apply the linear unit to all training examples, then compute for each weight according to

$$\Delta w_i = \eta \sum_d (t_d - o_d) x_{id}$$

- Update each  $w_i$  by adding  $\Delta w_i$  then repeat process.

# Gradient-Descent Algorithm

- Initialize each  $w_i$  to some small random value.
- Until the termination condition is met, do
  - Initialize each  $\Delta w_i$  to zero.
  - For each  $\langle x, t \rangle$  in training\_examples, do
    - Input the instance  $x$  to the unit and compute the output  $o$
    - For each linear unit weight  $w_i$ , do

$$\Delta w_i \leftarrow \Delta w_i + \eta(t - o)x_i$$

- For each linear unit weight  $w_i$ , do

$$w_i \leftarrow w_i + \Delta w_i$$

# Learning Rate

- Converges to a weight vector with minimum error
  - Since error surface contains a single global minimum
  - Training examples are linearly separable
  - Given a sufficiently small learning rate  $\eta$
- Large  $\eta$  runs the risk of overstepping the minimum in the error surface (i.e. not converging).
- We can also gradually reduce the value of  $\eta$  as the number of gradient descent steps grows.

# Stochastic Gradient Descent

- **Stochastic gradient descent** updates  $w_i$  for each training sample, normally randomly selected.

- Update weight as each sample is processed

$$\Delta w_i \leftarrow \Delta w_i + \eta(t - o)x_i$$

- Called Delta rule, LMS rule, Adaline rule, Widrow-Hoff rule.

# SGD Algorithm

- Initialize each  $w_i$  to some small random value.
- Until the termination condition is met, do
  - Initialize each  $\Delta w_i$  to zero.
  - For each  $\langle x, t \rangle$  in training examples, do
    - Input the instance  $x$  to the unit and compute the output  $o$
    - For each linear unit weight  $w_i$ , do

$$w_i \leftarrow w_i + \eta(t - o)x_i$$

# Exercise 2


- **Goal:** Try stochastic gradient on scikit-learn.
- **Instructions:**
  - Create datasets corresponding to ADD, OR and XOR.
  - Train `linear_model.SGDClassifier` with each dataset.
  - Check your results.
  - Extra: Try fitting a kernel before training (e.g. `RBFSampler`).

# Content


- Perceptrons
- Stochastic gradient descent algorithm
- Backpropagation & multi-layered networks
- Deep learning architectures & applications



# Multi-layered Neural Networks




# Decision Regions



# Differentiable Threshold Unit


- Highly non-linear multiple layer functions require functions that output a non-linear function whose output is a differentiable function; for example, a sigmoid.



# Sigmoid Function

- A logistic function:

$$y = 1/(1+e^{-x})$$



# Backpropagation Algorithm

- Until the termination condition is met, do
  - For each  $\langle x, t \rangle$  in training\_example, do
    - 1. Propagate the errors backwards through the network.
    - 2. For each network output unit  $k$ , calculate its error term  $\delta_k$

$$\delta_k \leftarrow o_k (1 - o_k) (t_k - o_k)$$

- 3. For each hidden unit  $h$ , calculate its error term  $\delta_h$

$$\delta_h \leftarrow o_h (1 - o_h) \sum_{k \in \text{outputs}} w_{kh} \delta_k$$

- 4. Update each network weight  $w_{ij}$  where

$$\Delta w_{ij} = \eta \delta_j x_{ji}$$

# Momentum

- Most common variation of BACKPROPAGATION is one that adds "momentum" to  $\Delta w_{ji}$ .
- Weight update on  $n^{\text{th}}$  iteration depends upon update that occurred in  $(n-1)^{\text{th}}$  iteration

$$\Delta w_{ji}(n) = \eta \delta_j x_{ji} + \alpha \Delta w_{ji}(n-1)$$

- where  $\alpha$  is the momentum.

# Multi-layered Networks

- Algorithm generalises to feedforward networks of arbitrary depth.
- For unit  $r$  in layer  $m$ ,  $\delta_r$  is computed from the values at next deeper layer  $m + 1$



$$\delta_r \leftarrow o_r (1 - o_r) \sum_{s \in \text{layer } m+1} w_{sr} \delta_s$$

# When to Terminate?

- Termination condition for backpropagation is not specified.
- Possible choice:
  - Continue training until error  $E$  falls below a pre-determined threshold.
  - This is a poor strategy because it can **overfit**.



# Generalisation vs Overfitting



# Stopping Criteria

- Termination conditions:
  - Fixed number of iterations through loop.
  - Once error on training examples fall below a threshold.
  - Error on a separate validation set meets a criterion.
- Choice of termination criterion important:
  - Too few iterations can fail to reduce error sufficiently.
  - Too many can lead to overfitting the training data.

# Exercise 3

- **Goal:** Train an iris classifier using MLPClassifier.
- **Instructions:**
  - Load iris dataset.
  - Build an MLPClassifier for iris classification.
  - Evaluate the performance of your neural network.

# Exercise 4

- **Goal:** Train digit classifier using MLPClassifier.
- **Instructions:**
  - Load digit dataset.
  - Build an MLPClassifier for handwritten digit recognition.
  - Evaluate the performance of your neural network.

# Exercise 5

- **Goal:** Train a sentiment analyser using MLPClassifier.
- **Instructions:**
  - Read data from provided text files for sentiment analysis problem.
  - Build an MLPClassifier for sentiment classification.
  - Evaluate the performance of your neural network.

# Content

- Perceptrons
- Stochastic gradient descent algorithm
- Backpropagation & multi-layered networks
- Deep learning architectures & applications

# Overview of Deep Learning


- Related to the way the brain *processes and communicates* information and patterns, to define relationships between *stimuli and responses*.
- More *feature learning* than *task-specific learning*.
- **Feature Learning**
  - Also called **representation learning**.
  - Allows systems to automatically discover representations needed for feature detection or classification from raw data.

# Common Architectures

- **Deep Neural Networks**
  - Artificial neural networks with many hidden layers.
- **Deep Belief Networks**
  - Bayesian networks with many layers representing hidden variables.
- **Deep Recurrent Networks**
  - Allow capture temporal behaviour (i.e. time sequence).
- **Convolutional Neural Networks**
  - Have a feedforward, convolutional layer as the key building block.




# Deep Neural Networks




# Deep Belief Networks

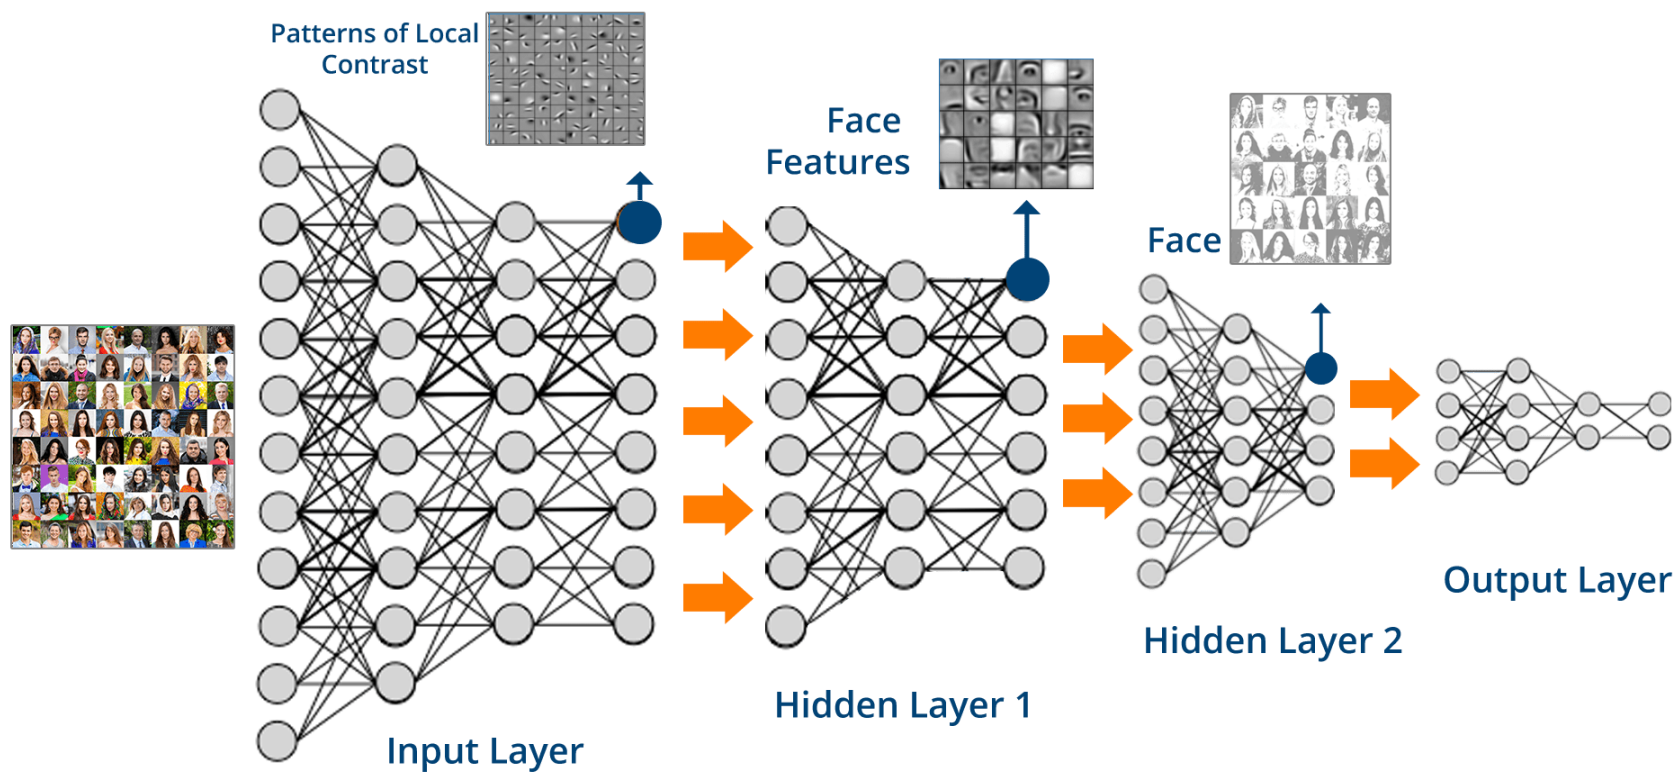
- Similar to deep neural networks, but connections are made between layers, NOT units within each layer.
- This allows a fast, layer-by-layer unsupervised training.
- Also can train a layer at a time.



# Long Short-term Memory (LSTM)



# Convolutional Neural Networks




# Components of a CNN

- **Convolutional Layers**
  - The core building block of CNN.
  - This layer is composed of a set of **kernels (filters)**.
- **Pooling Layers**
  - A component for sampling to reduce the size of representation.
  - For example: MaxPooling.
- **Activation Layers**
  - A component for applying an activation function, like standard ANN.
  - For example: ReLU units.


# Convolutional Layer

- Performs a **convolution** on the input data (e.g. an image).
- This is just like any convolution done in image processing to apply filters.
- The challenge is the choice of **learnable filters**, or **kernels**.



# Pooling Layers

- Performs **down-sampling** on the extracted features.
- We do not need the exact location of the feature, but only that *relative to other features*.
- Example: **MaxPooling**, where only the maximum is selected.



# Activation Layer

- The activation layer is analogous to the activation function used in backpropagation neural networks.
- It increases the non-linear properties of decision function.
- **Common activation functions:**
  - ReLU (Rectified Linear Units):  $\max(0, x)$
  - Logistic function, e.g.  $\text{sigmoid}(x)$
  - Other hyperbolic functions, e.g.  $\tanh(x)$